

L02: The Mathematical Building Blocks of Neural Networks

CSci 560 Deep Learning w/ Python (Chollet) Ch. 2

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

L02.1 A First Look at a Neural Network

CSci 560: Deep Learning w/ Python (Chollet) Ch. 2.1

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

L02.2 Data Representations for Neural Networks

CSci 560: Deep Learning w/ Python (Chollet) Ch. 2.2

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

L02.3 The Gears of Neural Networks: Tensor Operations

CSci 560: Deep Learning w/ Python (Chollet) Ch. 2.3

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

Matrix · Matrix Dot-Product

- x, y are input matrices of shape (a, b) and (b, c) respectively
- Each row of x performs an inner product with corresponding column of y
- So resulting element $z[i, j]$ is the inner product between row i of x and column j of y
- Resulting shape of z is (a, c)




Figure 1: Matrix dot-product box diagram. (Chollet, 2021, pg.43)



Vector Addition is Translation

- Adding a vector to a point will move the point by a fixed amount.
- Applied to a set of points, result is a translation.




Figure 2: 2D translation as vector addition. (Chollet, 2021, pg.45)



Matrix Multiplication (dot product) is Rotation

- A counterclockwise rotation by angle θ achieved by dot product.




Figure 3: 2D rotation (counterclockwise) as a dot product. (Chollet, 2021, pg.45)

Matrix Multiplication can also Scale

- Vertical and horizontal scaling also achieved by dot product with a suitable matrix.




Figure 4: 2D scaling as a dot product. (Chollet, 2021, pg.46)

Affine Transformation: Linear transform and translation

- Translation, Rotation, Scaling are all linear transformations.
- Can combine, for example $y = W \cdot x + b$ is a linear transformation plus a translation.




Figure 5: Affine transform in the plane. (Chollet, 2021, pg.46)

relu Nonlinear Activation Function

- The previous are just combinations of linear transformations.
- Multilayer NN made entirely of linear transformations can be reduced to a single linear transformation.
- Activation functions produce a nonlinear translation.




Figure 6: Affine transform followed by (nonlinear) 'relu' activation. (Chollet, 2021, pg.47)

A Geometric Interpretation of Deep Learning

- NN consist entirely of chains of tensor operations.
- Can be interpreted as geometric transformations of the input data into a new space/shape/manifold.
- Can interpret a NN as a very complex geometric transformation in a high-dimensional space, implemented via a series of simple linear + nonlinear transformations.



Figure 7: Uncrumpling a complicated manifold of data. (Chollet, 2021, pg.47)



L02.4 The Engine of Neural Networks: Gradient-based Optimization

CSci 560: Deep Learning w/ Python (Chollet) Ch. 2.4

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

1D Mini-Batch Stochastic Gradient Descent

- Stochastic because each batch of data is drawn at random.
- Need a reasonable value for the `learning_rate`.
 - too small, slow convergence
 - too large, diverges





Figure 8: SGD down a 1D loss curve (one learnable parameter) (Chollet, 2021, pg.53)

2D Mini-Batch Stochastic Gradient Descent



- Gradient descent will be used in highly dimensional spaces.
- Every weight coefficient in a neural network is a free dimension in the space.
 - Tens of thousands or even millions of parameters.
- Cannot visualize, and intuitions from low-dimension examples can sometimes be misleading.

Figure 9: Gradient descent down a 2D loss surface (two learnable parameters). (Chollet, 2021, pg.54)

SGD with Momentum / Advanced Optimization Methods




Figure 10: A local minimum and a global minimum.
(Chollet, 2021, pg.54)

- **Momentum** take into account previous weight updates rather than just looking at current gradient.
- SGD w/ momentum, Adagrad, RMSprop
- Momentum addresses two issues with SGD:
 - convergence speed and local minima.
 - If enough momentum won't get stuck in local minimum
 - Will adjust learning rate to to make sure convergence does not take too long.
- Implemented by moving each step based on both slope (current acceleration) and velocity.

Chaining Derivatives: The Backpropagation Algorithm


- **Backpropagation** is a way to use the derivatives of simple operations to easily compute the gradient of arbitrarily complex combinations of atomic operations.
- A NN consists of many tensor operations chained together, each of which has a simple known derivative.
- The chain rule states that
$$\text{grad}(y, x) == \text{grad}(y, x1) * \text{grad}(x1, x)$$
- Can be applied to any number of intermediate functions

$$\text{grad}(y, x) == (\text{grad}(y, x3) * \text{grad}(x3, x2) * \text{grad}(x2, x1) * \text{grad}(x1, x))$$

```
def fg(x):  
    x1 = g(x)  
    y = f(x1)  
    return y  
  
def fghj(x):  
    x1 = j(x)  
    x2 = h(x1)  
    x3 = g(x2)  
    y = f(x3)  
    return y
```



Computation Graphs and Automatic Differentiation



- Useful way to think of and perform backpropagation: use computation graphs.
- A **computation graph** is a directed acyclic graph of operations.
- Treat computation as data:
 - Graph can be built in TensorFlow from a GradientTape
 - Represents forward pass
 - Can reverse links to perform backpropagation backward pass.

Figure 11: The computation graph

Automatic Differentiation: Forward Pass

- Consider a simplified computation graph.
 - single layer all variables are scalar
 - scalar w for weights and b for bias
 - scalar x input
 - no activation function
 - loss is difference of result and y_{true}
- We want to update w and b in a way that will minimize loss_val .
 - want to compute $\text{grad}(\text{loss_val}, b)$ and $\text{grad}(\text{loss_val}, w)$.
- Propagate these values to all nodes in the graph from top to bottom, until we reach loss_val .
 - This is the **forward pass**




Figure 12: Running a forward pass.
(Chollet, 2021, pg.58)

Automatic Differentiation: Backward Pass / Backpropagation

- “reverse” the graph
 - For each edge from A to B, create opposite edge from B to A
 - backward graph represents the **backward pass**, a concrete implementation of performing **backpropagation**
 - You can obtain the derivative of a node with respect to another node by *multiplying the derivatives for each edge along the path linking the two nodes*.




Figure 13: Running a forward pass. (Chollet, 2021, pg.58)



L02.5 Looking Back at our First Neural Network Example

CSci 560: Deep Learning w/ Python (Chollet) Ch. 2.5

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2025

Chollet, F. (2021). *Deep learning with python* (second). [Manning](#).